

Tool-use summary

Santheesh S | Track B | Product 1 (DICM Core)

Total time on this submission: 3 hours 50 minutes (230 min).

Tools I used

1. Cursor (Composer) - AI assistant in the IDE

What I used

How I checked output: Re-ran `python product1\verify_product1.py` (PRODUCT 1 PASS) and `python -m unittest discover -s tests -v` after changes. Compared behaviour to `DATA_DICTIONARY.md`; routed conflicting pending + received_at to uncertain.

2. Python 3.10+ (stdlib only; no paid APIs)

What I used

How I checked output: Verify on unchanged Mystri data/; baseline 13 vs rules 5 locked in `test_pack_snapshot_baseline_thirteen_rules_five`.

3. Manual analysis

What I used

DECISION.r

How I checked output: Cross-checked CSVs and test output; did not treat the owner 8 hours/week claim as measured fact.

4. Web / vendor documentation

What I used

How I checked output: Recorded in `SOURCES.md` with access dates, supported claims, and limitations.

I did not use paid LLM APIs, deployment platforms, or live customer messaging for this submission.

One concrete issue I caught

The assistant used an invalid dataclass update pattern in queue-related code. Unit tests failed. I replaced it with `dataclasses.replace`, re-ran the full test suite, and confirmed PRODUCT 1 PASS before treating the work as submission-ready.

AI vs my judgment

AI: boilerplate, test stubs, and doc drafts.

Me: final rule dispositions, numbers in `DECISION.md`, Product 1 scope, and relying on results only after verify and tests pass.

Contact: sivasandy509@gmail.com