

Tool-Use Summary

Santheesh S · Mystri Track B · Product 1 (DICM Core) · 3 h 50 min (230 min)

sivasandy509@gmail.com

1. Why I document tool use

Mystri requires transparency on how I produced this submission. I treat **AI as an accelerator**, not the author of record. My evidence is what runs on the Mystri pack (**python product1\verify_product1.py** → **PRODUCT 1 PASS**, 39 tests) and what I wrote in **DECISION.md** and **SOURCES.md**.

2. My working theory (AI on a timed assessment)

- **Human owns the decision; AI owns drafts.** Recommendation, rule semantics, and scope are mine after review.
- **Verify before trust.** AI code must pass the same gate a reviewer uses: verify + tests on unchanged **data/**.
- **Collection ≠ policy.** File-request tools help uploads; they do not replace 48h gaps, opt-outs, or row reconciliation—this shaped DICM.
- **Fail closed on ambiguity.** Conflicting fields (e.g. **R018**) → **uncertain**, not auto-contact.
- **State limits.** I do not claim live savings, LLM necessity, or the owner's 8 h/week as measured fact.

3. Tools I used (mandatory disclosure)

3.1 Cursor — Composer (AI coding assistant)

What: IDE agent in Cursor; primary AI pair-programmer for this repo.

Model/settings: Composer as provided by Cursor; no separate paid LLM API for graded runs.

Used for: Scaffolding **queue_engine.py**, parts of **integrated_model.py**, test stubs, doc structure.

Not for: Final rule decisions, **DECISION.md** numbers, or success without **PRODUCT 1 PASS**.

Checked by: **verify_product1.py** + full unittest; compared to **DATA_DICTIONARY.md**.

3.2 Python 3.10+ (stdlib only)

Used for: Loader, experiment, verify gate, 39 tests, dry-run JSON/log outputs.

Checked by: **PRODUCT 1 PASS**; baseline **13** vs rules **5** in **test_pack_snapshot_baseline_thirteen_rules_five**.

Theory: Fixed snapshot (**scenario.json**) gives reproducible proof without credentials.

3.3 Manual analysis

Used for: Deduped **event_id**; naive vs rules counts in **DECISION.md**.

Checked by: CSV + test cross-check.

Rejected: Owner “8 h/week” as unmeasured vs logged minutes.

3.4 Web / vendor documentation

Used for: OneDrive/Dropbox file-request, Microsoft Learn share-files, Jotform (limited fit).

Checked by: **SOURCES.md** with dates, claims, limits.

Theory: Tested “buy tool fixes chasing”—docs support collection, not Daybreak follow-up policy.

4. What I did not use

Paid LLM APIs for grading runs; deployment; live email/SMS/WhatsApp; customer interviews; edits to Mystri **data/** for graded results.

5. Concrete example — issue I caught

Composer used an **invalid dataclass update** → **unit tests failed**. I **rejected** that output, fixed with **dataclasses.replace**, re-ran all tests and verify until **PRODUCT 1 PASS**. AI is a draft until that gate passes.

6. AI vs my judgment

AI: boilerplate, test skeletons, first-pass docs. | **Me:** rules, evidence, scope, pass/fail criteria. | **Reject rule:** fail verify/tests → fix or revert.

Repo: github.com/santheesh15/-mystri-assessment · **Branch:** [cursor/track-b-submission-edb7](#) · **Also in repo:** [track-b/docs/TOOL_USE_SUMMARY.md](#)